

# COL774- MACHINE LEARNING


## ASSIGNMENT 4

DIWAKAR PRAJAPATI - 2018CS10330

UTSAV DEEP - 2018CS10396

### DATA SET

The dataset contains 7 types of facial expressions on grayscale. A few of the samples with the class labels are given below.



## NON COMPETITIVE PART

We have used Pytorch Library for the neural network.

### A. VANILLA NEURAL NETWORK

No of Epochs	Learning Rate	Weight decay	Activation function	Training Cost	Training Time	Train Accuracy	Test Accuracy
200	0.005	0.0001	ReLU	0.033493	4min 10secs	99.77%	39.45%
200	0.005	0.0001	Sigmoid	0.728391	4min 27secs	82.71%	38.26%
200	0.005	0.0001	Tanh	0.052944	4 min 28secs	99.73%	36.72%


- > The specifications of the layers were as given in the problem statement.
- > We tried 3 different activation functions - ReLU, Sigmoid, Tanh. Cross entropy loss is used to train the data.
- > From the results we can see that among various activation functions, ReLU activation gives the best accuracy. Next, best is by Sigmoid and last by Tanh.
- > Training accuracy of both ReLU and Tanh were very high ~ 100% but the training accuracy for sigmoid was very low.
- > Training time was almost the same for all parts.
- > ReLU function has been used for final submission. In all the subsequent parts also ReLU will only be used as an activation function unless specified explicitly.

### B. FEATURE ENGINEERING

The filters were implemented using the skimage library for image processing.

I. Applying Gabor filters :

Applying different Gabor filters to an image resulted in different outputs as shown below:



Parameters for first filter: frequency=0.9, theta = 0, sigma\_x=0.5, sigma\_y=0.5

Parameters for second filter: frequency=0.9, theta = 0, sigma\_x=0.8, sigma\_y=0.8

Parameters for third filter: frequency=0.9, theta = 0, sigma\_x=0.9, sigma\_y=0.9

> We passed each image through each of the above filters and then trained the model.

> We did this to reduce the overfitting in the model. Thus each input image would transform into three images from the Gabor filters.

Accuracy values after passing through Gabor filter:

Number of epochs: 200


Learning rate: 0.005

Training accuracy: 99.82 %

Test accuracy: 40.12 %

ii. Histogram of Oriented Gradients Filter( HOG filter):

Applying hog filter to an image resulted in the following



The parameters used were:

`orientations=16, pixels_per_cell=(3, 3), cells_per_block=(2, 2), visualize=True, multichannel=False`

> The hog filter was taking a lot of time to train.

Accuracy values after passing through HOG filter:

Number of epochs: 200

Learning rate: 0.005

Training accuracy: 99.80 %

Test accuracy: 39.94 %

> The time to train has increased as time is taken to apply filter on each image. Also hog filter was taking a lot more time compared to gabor filter.

> We observe that using the image filters has helped improve the accuracy, however that increase is only little and not very huge.

> The best performing feature engineering was by Gabor filter and that has been used in the final submission.

## C. CONVOLUTIONAL NEURAL NETWORK

Optimizer	No of Epochs	Learning Rate	Weight Decay	Training Cost	Train Accuracy	Test Accuracy
SGD	200	0.001	0.0001	0.102990	97.50%	41.59%
Adam	200	0.001	0.0001	0.103110	97.48%	40.46%

> We used SGD optimiser on the following network structure:

```
digit_model_cnn = torch.nn.Sequential(
    torch.nn.Conv2d(kernel_size=(3,3), stride=3, padding=0, in_channels=1, out_channels=64),
    torch.nn.BatchNorm2d(64),
    torch.nn.ReLU(),
    torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=0),
    torch.nn.Conv2d(kernel_size=(2,2), stride=2, padding=0, in_channels=64, out_channels=128),
    torch.nn.BatchNorm2d(128),
    torch.nn.ReLU(),
    torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=0),
    torch.nn.Flatten(start_dim=1),
    torch.nn.Linear(128,32),
    torch.nn.ReLU(),
    torch.nn.BatchNorm1d(32),
    torch.nn.Linear(32,10)
)
digit_model_cnn
```

> Inference time taken by Neural Networks is lesser than that of CNN.

> However, f1-score is almost the same and there is not a lot of difference.

> Best accuracy after a lot of submissions was seen when we first applied the Gabor filter on the image and then apply the CNN network defined in question 3 which resulted in about 42 % accuracy.

## COMPETITIVE PART

- > We used CNN for this part without a gabor filter or hog filter.
- > We have augmented the data before training.
- > The architecture we used for this part was as follows:

```
model = torch.nn.Sequential(  
    torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),  
    torch.nn.ReLU(inplace=True),  
    torch.nn.BatchNorm2d(32),  
    torch.nn.MaxPool2d(kernel_size=2, stride=2),  
    torch.nn.Dropout(p=0.25),  
    torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),  
    torch.nn.ReLU(inplace=True),  
    torch.nn.BatchNorm2d(64),  
    torch.nn.MaxPool2d(kernel_size=2, stride=2),  
    torch.nn.Dropout(p=0.25),  
    torch.nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),  
    torch.nn.ReLU(inplace=True),  
    torch.nn.BatchNorm2d(128),  
    torch.nn.MaxPool2d(kernel_size=2, stride=2),  
    torch.nn.Dropout(p=0.25),  
    torch.nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1),  
    torch.nn.ReLU(inplace=True),  
    torch.nn.BatchNorm2d(128),  
)
```

```
torch.nn.MaxPool2d(kernel_size=2, stride=2),  
torch.nn.Dropout(p=0.25),  
torch.nn.Flatten(start_dim=1),  
torch.nn.Linear(1152, 512),  
torch.nn.ReLU(inplace=True),  
torch.nn.Dropout(),  
torch.nn.Linear(512, 256),  
torch.nn.ReLU(inplace=True),  
torch.nn.Dropout(),  
torch.nn.Linear(256, 7),  
)
```

> Before passing through the filter the image was augmented with the following transforms:

- 1) 45-degree anti-clockwise rotation.
- 2) Flip up-down.
- 3) Flip left-right.
- 4) Random noise.

The final accuracy we obtained was:

54% on the test data

85% on training data.

THANK YOU!